

Sparse Recovery Above the Phase Transition: A Learned Operator-Aware Approach

Di An
Johns Hopkins University
Email: dan5@jhu.edu

Dylan Poppert
Johns Hopkins University
Email: dpopper2@jhu.edu

Taewoon Choi
Johns Hopkins University
Email: tchoi10@jh.edu

Trac D. Tran
Johns Hopkins University
Email: trac@jhu.edu

Abstract—We characterize the sparsity regime where learned sparse recovery has genuine headroom over classical methods. Below the ℓ_1 phase transition ($k=25$ in $n=256$, $m=128$), CoSaMP [2] achieves exact recovery on both partial Fourier and Gaussian operators, while LISTA [8] trained on Fourier and tested zero-shot on Gaussian plateaus at NRMSE 0.33. At $k=55$ (just above the Donoho–Tanner weak threshold [6]), polynomial-time greedy methods fail (CoSaMP=0.55, OMP=0.90 [3] NRMSE) while oracle-support least squares remains exact. A coordinate-wise learned support detector with operator-aware features beats coherence-only baselines (25–34% relative NRMSE) but does not match CoSaMP. The over-sparse regime is the open setting where learned recovery has a defensible role, with two structural advantages — adaptive cardinality and amortized signal-prior inference — independent of regime. On Gaussian compressible signals the same detector, trained on mixed-tail-amplitude signals, beats CoSaMP by 0.074–0.097 NRMSE (8–10% relative) at off-support tail amplitudes 0.3–0.4 (where CoSaMP’s NRMSE exceeds 1); the gap is structural, holding in 15 of 15 runs across 5 operator draws $\times$ 3 initializations. The broader aim is a learning approach that beats CoSaMP across all sparsity regimes — the compressibility win is the first concrete instance.

I. BELOW THE TRANSITION: CLASSICAL SPARSE RECOVERY WINS

We consider the sparse linear inverse problem $y = Ax + \varepsilon$ [1] with $A \in \mathbb{R}^{m \times n}$ fixed at deployment and $\|x\|_0 = k$. Classical baselines (CoSaMP [2], OMP [3], AMP [4]) and learned-recovery work (LISTA [8] and variants [9]–[11]) are typically benchmarked well below the ℓ_1 phase transition [6]. At $n=256$, $m=128$, $k=25$ (noiseless), CoSaMP achieves NRMSE ≈ 0 on both partial Fourier ($\mu=0.75$) and Gaussian ($\mu=0.37$) matrices. By contrast, LISTA trained on Fourier and tested zero-shot on Gaussian achieves NRMSE 0.33; the best parameter-free support-then-amplitude baseline achieves 0.26. The learning “advantage” here is illusory: CoSaMP [2] dominates by an order of magnitude in this strict-sparse regime. (The picture changes once signals become compressible rather than exactly k -sparse: see Section V.)

II. ABOVE THE TRANSITION: LEARNING APPROACH HAS THE UPPER HAND

The landscape inverts above the phase transition. Table I reports NRMSE on Gaussian at $k=55$ ($k/n=0.21$, $k/m=0.43$, just above the Donoho–Tanner weak threshold $\rho_W \approx 0.4$): signals remain sparse, but classical guarantees no longer hold. All polynomial-time greedy methods fail by 0.55–0.90 NRMSE

TABLE I
GAUSSIAN OPERATOR ($n=256$, $m=128$, 500 TEST SIGNALS). NRMSE AT $k=55$, ABOVE THE ℓ_1 PHASE TRANSITION.

Method	NRMSE
Naive top- k on $ A^\top y $ + LS	0.79
OMP	0.90
CoSaMP	0.55
Learned detector (ours)	0.59
Oracle support + LS	0.00

while oracle support recovery remains exact: the recovery problem is well-posed but combinatorial — the regime in which learning approaches have headroom (Section V).

III. OPERATOR-AWARE SUPPORT DETECTOR

For each coordinate j we build a feature vector $\phi_j = [|x_j^{(T)}|, A_{:,j}^\top(y - Ax^{(T)}), A_{:,j}^\top y, (A^\top A)_{jj}, \mu_j]$ with $x^{(T)}$ a $T=30$ -step ISTA [5] iterate and $\mu_j = \max_{i \neq j} |A_{:,i}^\top A_{:,j}|$. The residual correlation encodes whether j can still explain the measurement residual; the Gram diagonal and coherence describe local identifiability under the deployment operator. A coordinate-wise MLP (4,609 params) maps each ϕ_j to a support probability; the top- k scoring indices form the predicted support, with amplitudes recovered by support-restricted least squares.

Trained on 4,000 $k=55$ -sparse signals and evaluated on 500 held-out signals, the detector achieves NRMSE 0.59 vs 0.79 for naive top- k and 0.90 for OMP; it does not match CoSaMP’s 0.55. Test IoU plateaus at 0.51 by epoch 10 — an information ceiling, not an optimization one.

IV. DISCUSSION

The CoSaMP gap is structural: CoSaMP performs implicit set-level reasoning (merge $2k$, refit, prune), while a coordinate-wise detector scores each index from its own feature vector with no view of competing candidates. The IoU plateau at 0.51 pins this as a representational ceiling, not an optimization one. Two further structural advantages of learned detectors hold *independently* of the recovery regime. First, CoSaMP requires k as input, while a learned threshold-based detector can estimate cardinality from data — real signals are not exactly k -sparse and the true k is rarely known. Fig. 1 (right column) shows the qualitative trend; Section V turns it into a quantitative win. Second, classical methods are per-instance




Fig. 1. Stress tests on partial Fourier ($\mu=0.75$, top) and Gaussian ($\mu=0.37$, bottom) operators at $n=256$, $m=128$. **Left:** sparsity sweep—CoSaMP and OMP achieve exact recovery below the phase transition but degrade sharply past $k \approx 40$, while oracle support recovery remains exact throughout. **Center:** noise sweep at $k=25$ —all methods track the oracle closely above 15 dB. **Right:** approximate-sparsity sweep—as off-support tail amplitude grows, oracle performance itself degrades and the greedy-naïve gap narrows, exposing CoSaMP’s reliance on exact k -sparsity.

and oblivious to signal priors, while learning amortizes prior structure across the training distribution — a gain that should compound for signals with exploitable structure (block-sparse, transform-domain compressible).

V. BEYOND COSAMP: AN EMPIRICAL WIN UNDER COMPRESSIBILITY

The structural argument — that CoSaMP’s strict k -sparsity assumption becomes the bottleneck once signals are compressible rather than exactly sparse — is backed by a direct empirical win. Training the same coordinate-wise detector on Gaussian compressible signals [7] with random off-support tail amplitude $\sim \text{Unif}[0.1, 0.4]$ ($k=25$, $m=128$, labels = top- k of $|x|$) and evaluating at fixed tail amplitudes (500 test signals each):

tail amplitude	0.1	0.2	0.3	0.4
CoSaMP	0.28	0.63	0.88	1.01
Learned (ours)	0.37	0.62	0.80	0.91
Δ (gain)	-0.08	+0.01	+0.07	+0.10

NRMSE; bold marks the lower (better) value per column. $\Delta = \text{CoSaMP} - \text{Learned}$, so positive favors the learned detector.

Each cell reports the mean across 5 operator seeds $\times$ 3 model initializations (15 runs); the wins at tail 0.3 and 0.4 hold in 15 of 15 runs with cross-run std ≤ 0.003 NRMSE (paired-bootstrap 95% CI: $[+0.073, +0.076]$ at tail 0.3, $[+0.096, +0.099]$ at tail 0.4). The result is structural, not a single-shot artifact: CoSaMP wins at low tail (its strict- k assumption nearly holds) but collapses past tail 0.2; at

0.4 its NRMSE exceeds 1, worse than predicting zero. The learned detector crosses CoSaMP at tail 0.2 and pulls 8–10% ahead at 0.3–0.4, the only method that stays below the 1.0 floor. The mechanism is exactly what the structural argument predicted: top- k of $|x|$ now includes the largest tail entries, which CoSaMP — locked to the original spike support — cannot select. The low-tail loss is also structural: widening training to $\text{Unif}[0, 0.4]$ does not close it, indicating that per-coordinate features cannot match CoSaMP’s set-level merge-and-refit at the strict-sparse end. Note also that on strict-sparse $k=55$ Gaussian the detector already matches CoSaMP’s support IoU (0.51 vs 0.53), so the residual NRMSE gap there is amplitude- not support-bound — a clue Section VI follows up.

VI. PRELIMINARY: ATTENTION AS A SET-LEVEL MECHANISM

The IoU plateau at 0.51 motivated a single-block remedy: replace the per-coordinate MLP with a small transformer encoder [12] ($n=256$ coordinate tokens, embedding dim 24, 4 heads, 2 layers, $\sim 10k$ params), so each index can see its competitors. We inject operator awareness through an additive Gram-matrix attention bias: for each head h , the attention logits gain $\alpha_h \cdot |G_{ij}|$ before the softmax, where $G=A^T A$ with the diagonal zeroed and α_h is a learnable scalar initialized at 0. We call this GAST — the Gram-Aware Support Transformer. Two preliminary single-seed findings, paired against the MLP on the same operator and training data:




Fig. 2. Per-head, per-layer Gram-bias scaling α_h over 80 training epochs. Solid: layer 0; dashed: layer 1. Both regimes drive α to magnitudes of ± 2.5 with a mix of positive (attend *more* to coherent neighbors) and negative (attend *away*) heads. The mechanism engaged equivalently in both regimes, yet only compressible translates that engagement into a test-time win.

(i) **Compressibility extension.** On compressible $k=25$ Gaussian (Unif[0.1, 0.4] training), GAST extends the MLP’s CoSaMP win in every cell of the §V table: at tail= 0.4, GAST beats CoSaMP by 0.106 NRMSE versus MLP’s 0.097; at tail= 0.3, +0.080 versus +0.075. Modest but consistently in the right direction.

(ii) **Strict-sparse no-improvement.** On strict $k=55$ Gaussian, GAST and MLP land within 0.001 NRMSE and 0.001 IoU of each other — both plateau, both lose to CoSaMP by ~ 0.034 . The 0.51 IoU ceiling holds.

The negative result is informative because α *did* engage (Fig. 2): magnitudes reach ± 2.5 in both regimes, with positive heads (focus on competing candidates) and negative heads (suppress confusable neighbors); trajectories are essentially indistinguishable across regimes, yet only compressible yields a test gain. The IoU plateau is therefore *not* a per-coord-vs-set-level distinction — set-level reasoning was added with maximum expressiveness, α used it, and test performance still did not move. The remaining bottleneck on strict-sparse is more likely either an information ceiling in the 5-feature interface, or the absence of iterative algorithmic structure (CoSaMP’s merge-refit-prune loop runs against the *current* residual; neither MLP nor a single-block GAST does).

VII. OUTLOOK

The preliminary GAST result reframes the next-architecture question: single-block attention is necessary but not sufficient at strict-sparse $k=55$; the natural escalation is iterative algorithmic structure — a learnable analogue of CoSaMP’s merge-refit-prune loop with attention modules at each step. Two experiments would convert the regime-independent advantages from claim to evidence: (i) training on block- or cluster-sparse priors to measure the amortization gain; and (ii) replacing top- k with a learned threshold to test adaptive cardinality directly. Together — iterative attention, learned threshold, structured priors — these directions target all three of CoSaMP’s structural limitations (set-level reasoning, prior-blindness, oracle- k dependence), aiming at a learned operator-aware framework that uniformly outperforms CoSaMP regardless of sparsity regime.

REFERENCES

- [1] D. L. Donoho, “Compressed sensing,” *IEEE Trans. Inf. Theory*, vol. 52, no. 4, pp. 1289–1306, 2006.
- [2] D. Needell and J. A. Tropp, “CoSaMP: Iterative signal recovery from incomplete and inaccurate samples,” *Applied and Computational Harmonic Analysis*, vol. 26, no. 3, pp. 301–321, 2009.
- [3] J. A. Tropp and A. C. Gilbert, “Signal recovery from random measurements via orthogonal matching pursuit,” *IEEE Trans. Inf. Theory*, vol. 53, no. 12, pp. 4655–4666, 2007.
- [4] D. L. Donoho, A. Maleki, and A. Montanari, “Message-passing algorithms for compressed sensing,” *Proc. Nat. Acad. Sci.*, vol. 106, no. 45, pp. 18914–18919, 2009.
- [5] I. Daubechies, M. DeFrise, and C. De Mol, “An iterative thresholding algorithm for linear inverse problems with a sparsity constraint,” *Comm. Pure Appl. Math.*, vol. 57, no. 11, pp. 1413–1457, 2004.
- [6] D. L. Donoho and J. Tanner, “Counting faces of randomly projected hypercubes and orthants, with applications,” *Discrete & Computational Geometry*, vol. 43, no. 3, pp. 522–541, 2010.
- [7] A. Cohen, W. Dahmen, and R. DeVore, “Compressed sensing and best k -term approximation,” *J. Amer. Math. Soc.*, vol. 22, no. 1, pp. 211–231, 2009.
- [8] K. Gregor and Y. LeCun, “Learning fast approximations of sparse coding,” in *Proc. Int. Conf. on Machine Learning (ICML)*, 2010, pp. 399–406.
- [9] X. Chen, J. Liu, Z. Wang, and W. Yin, “Theoretical linear convergence of unfolded ISTA and its practical weights and thresholds,” in *Proc. Adv. Neural Inf. Process. Syst. (NeurIPS)*, 2018, pp. 9061–9071.
- [10] J. Liu, X. Chen, Z. Wang, and W. Yin, “ALISTA: Analytic weights are as good as learned weights in LISTA,” in *Proc. Int. Conf. on Learning Representations (ICLR)*, 2019.
- [11] M. Borgerding, P. Schniter, and S. Rangan, “AMP-inspired deep networks for sparse linear inverse problems,” *IEEE Trans. Signal Processing*, vol. 65, no. 16, pp. 4293–4308, 2017.
- [12] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin, “Attention is all you need,” in *Proc. Adv. Neural Inf. Process. Syst. (NeurIPS)*, 2017, pp. 5998–6008.